

Alexis J. Ross

alexisr@allenai.org
<https://alexisjihyeross.github.io/>

EDUCATION	Harvard University , Cambridge, Massachusetts A.B. with Honors in Computer Science and Philosophy (Joint)	Sept 2016 – May 2020
RESEARCH EXPERIENCE	Allen Institute for Artificial Intelligence (AI2), AllenNLP Predoctoral Researcher <i>Mentors:</i> Ana Marasović and Matthew E. Peters	July 2020 – Present
	Harvard University, Computer Science and Philosophy <i>Mentors:</i> Himabindu Lakkaraju and Bernhard Nickel <i>Senior Thesis:</i> Using Linear Approximations to Explain Complex, Blackbox Classifiers Awarded Thomas T. Hoopes Prize for “outstanding scholarly work” by undergraduates.	Sept 2019 – May 2020
	Microsoft Research Research Intern, Project Hanover <i>Mentor:</i> Tristan Naumann <i>Project:</i> Machine Reading for Precision Medicine	May 2019 – Aug 2019
	Harvard University, NLP Group Research Assistant <i>Mentors:</i> Yonatan Belinkov and Alexander Rush <i>Project:</i> Probing for Neurons that Encode and Exploit Negation in BERT	March 2019 – May 2019
	Johns Hopkins, Center for Language and Speech Processing Undergraduate Researcher, JSALT Workshop <i>Mentor:</i> Ellie Pavlick <i>Project:</i> General-Purpose Sentence Representation Learning	June 2018 – Aug 2018
	Harvard University, Cognitive Neuropsychology Laboratory Research Assistant <i>Mentors:</i> Gilles Vannuscorps and Alfonso Caramazza <i>Project:</i> Object Orientation Perception	Feb 2017 – Aug 2017
PREPRINTS	[1] Tailor: Generating and Perturbing Text with Semantic Controls. Alexis Ross* , Tongshuang Wu*, Hao Peng, Matthew Peters, Matt Gardner. In submission, 2021.	
PUBLICATIONS	[2] Learning Models for Actionable Recourse. Alexis Ross , Himabindu Lakkaraju, and Osbert Bastani. <i>NeurIPS</i> , 2021.	
	[3] Competency Problems: On Finding and Removing Artifacts in Language Data. Matt Gardner*, William Merrill*, Jesse Dodge, Matthew E. Peters, Alexis Ross , Sameer Singh, Noah Smith. <i>EMNLP</i> , 2021.	
	[4] Explaining NLP Models via Minimal Contrastive Editing (MiCE). Alexis Ross , Ana Marasović, and Matthew E. Peters. <i>Findings of ACL 2021</i> .	
	[5] How well do NLI models capture verb veridicality? Alexis Ross and Ellie Pavlick. <i>EMNLP 2019</i> . Oral presentation .	
	[6] Probing what different NLP tasks teach machines about function word comprehension. Najoung Kim, Roma Patel, Adam Poliak, Alex Wang, Patrick Xia, R Thomas McCoy, Ian Tenney, Alexis Ross , Tal Linzen, Benjamin Van Durme, Samuel R Bowman, and Ellie Pavlick. <i>*SEM 2019</i> . Best Paper Award .	

TEACHING
EXPERIENCE

CS 121: Introduction to the Theory of Computation

Sept 2018 – Dec 2018

Teaching Fellow, Harvard School of Engineering and Applied Sciences

Instructor: Boaz Barak

Awarded a Certificate of Distinction in Teaching for a student evaluation of 4.8/5.0.