

# Alexis J. Ross

[alexisr@allenai.org](mailto:alexisr@allenai.org)  
<https://alexisjihyeross.github.io/>

---

EDUCATION	<b>Harvard University</b> , Cambridge, Massachusetts A.B. with Honors in Computer Science and Philosophy (Joint)	Sept 2016 – May 2020
RESEARCH EXPERIENCE	<b>Allen Institute for AI (AI2), AllenNLP</b> Predoctoral Researcher <i>Mentors:</i> Ana Marasović and Matthew E. Peters	July 2020 – Present
	<b>Harvard University, Computer Science and Philosophy</b> <i>Mentors:</i> Himabindu Lakkaraju and Bernhard Nickel <i>Senior Thesis:</i> <a href="#">Using Linear Approximations to Explain Complex, Blackbox Classifiers</a>	Sept 2019 – May 2020
	<b>Microsoft Research</b> Research Intern, Project Hanover <i>Mentor:</i> Tristan Naumann <i>Project:</i> Machine Reading for Precision Medicine	May 2019 – Aug 2019
	<b>Harvard University, NLP Group</b> Research Assistant <i>Mentors:</i> Yonatan Belinkov and Alexander Rush <i>Project:</i> Probing for Neurons that Encode and Exploit Negation in BERT	March 2019 – May 2019
	<b>Johns Hopkins, Center for Language and Speech Processing</b> Undergraduate Researcher, JSALT Workshop <i>Mentor:</i> Ellie Pavlick <i>Project:</i> General-Purpose Sentence Representation Learning	June 2018 – Aug 2018
	<b>Harvard University, Cognitive Neuropsychology Laboratory</b> Research Assistant <i>Mentors:</i> Gilles Vannuscorps and Alfonso Caramazza <i>Project:</i> Object Orientation Perception	Feb 2017 – Aug 2017
PREPRINTS	[1] <a href="#">Tailor: Generating and Perturbing Text with Semantic Controls.</a> <b>Alexis Ross*</b> , Tongshuang Wu*, Hao Peng, Matthew Peters, Matt Gardner. In submission, 2021.	
PUBLICATIONS	[2] <a href="#">Learning Models for Actionable Recourse.</a> <b>Alexis Ross</b> , Himabindu Lakkaraju, and Osbert Bastani. <i>NeurIPS</i> , 2021.	
	[3] <a href="#">Competency Problems: On Finding and Removing Artifacts in Language Data.</a> Matt Gardner*, William Merrill*, Jesse Dodge, Matthew E. Peters, <b>Alexis Ross</b> , Sameer Singh, Noah Smith. <i>EMNLP</i> , 2021.	
	[4] <a href="#">Explaining NLP Models via Minimal Contrastive Editing (MiCE).</a> <b>Alexis Ross</b> , Ana Marasović, and Matthew E. Peters. <i>Findings of ACL 2021</i> .	
	[5] <a href="#">How well do NLI models capture verb veridicality?</a> <b>Alexis Ross</b> and Ellie Pavlick. <i>EMNLP 2019</i> . <b>Oral presentation.</b>	
	[6] <a href="#">Probing what different NLP tasks teach machines about function word comprehension.</a> Najoung Kim, Roma Patel, Adam Poliak, Alex Wang, Patrick Xia, R Thomas McCoy, Ian Tenney, <b>Alexis Ross</b> , Tal Linzen, Benjamin Van Durme, Samuel R Bowman, and Ellie Pavlick. <i>*SEM 2019</i> . <b>Best Paper Award.</b>	

SELECTED HONORS	2020	<b>Hoopes Prize</b> “for extraordinary undergraduate research,” awarded for senior thesis	
	2019	<b>EMNLP Student Travel Scholarship</b>	
	2018	<b>Certificate of Distinction in Teaching</b> for a student evaluation of 4.8/5.0	
SERVICE	2021	<b>Co-organizer</b> and frequent host, NLP Highlights Podcast, over 9000 views in October	
	2021	<b>Co-organizer</b> , AllenNLP Hacks, a hackathon for minoritized students Presented workshop: “My First NLP Project”	
	2021	<b>Secondary Reviewer</b> , ACL	
	2019	<b>Mentor</b> , Harvard Women in Computer Science	
TEACHING EXPERIENCE	<b>CS 121: Introduction to the Theory of Computation</b>		Sept 2018 – Dec 2018
	Teaching Fellow, Harvard School of Engineering and Applied Sciences <i>Instructor:</i> Boaz Barak		
OTHER EXPERIENCE	<b>Solo, Chamber, and Orchestral Violinist</b>		2002 – 2020
	2016 – 2018	Violinist in the Harvard-Radcliffe Orchestra	
	2016 – 2020	Chamber Music Performance at Harvard. Coaches: Parker Quartet	
	2014 – 2016	Juilliard Pre-College, Diploma in Violin Performance. Teacher: Ann Setzer	
	2013 – 2015	Summers at the Meadowmount School of Music	
	<b>Mental Health Advocacy and Counseling, Harvard University</b>		2017 – 2020
	2017 – 2020	Events Co-Chair, Student Mental Health Liaisons	
	2018 – 2020	Peer Mental Health Counselor, <a href="#">Room 13</a>	